

Computer Vision

Kenneth.Dawson-Howe@scss.tcd.ie

School of Computer Science and Statistics,
Trinity College, University of Dublin. Ireland .

Introduction

Based on *A Practical Introduction to Computer Vision with
OpenCV* by Kenneth Dawson-Howe © Wiley & Sons Inc. 2014

Slide 1

1

Module contents

Week	Beginning	Folder	Topic	Format	Mini Tests (10%)	Live Classes	Practical Component
1	12/09/2022	Introduction	What is Computer Vision	Live		Lecture Tuesday @2pm	
		Introduction	Module Organisation	Live		Lecture Wednesday @1pm	Install OpenCV & TIPS.
		Images	Camera Model & Digitisation	Pre-recorded			Experiment with TIPS.
2	19/09/2022	Images	Colour Image models	Pre-recorded			Assignment 1 (30%), Part 1
		Images	Noise & Smoothing	Pre-recorded	Images before Tuesday @12noon	Q&A Session Tuesday @2pm	
		Binary	Thresholding	Pre-recorded		Tutorial Wednesday @1pm	
3	26/09/2022	Binary	Cleaning Binary Images	Pre-recorded	Binary before Tuesday @12noon		
		Regions	Connectivity & Connected Component Analysis	Pre-recorded			Assignment 1, Part 2
		Regions	kMeans clustering	Pre-recorded	Regions before Tuesday @12noon	Q&A Session Tuesday @2pm	
4	03/10/2022	Geometric	Transformations	Pre-recorded		Tutorial Wednesday @1pm	
		Geometric	Interpolation	Pre-recorded	Geometric before Tuesday @12noon		Assignment 1, Part 3
		Histograms	1D & 3D histograms	Pre-recorded			
5	10/10/2022	Histograms	Equalisation, Comparison & Back Projection	Pre-recorded	Histograms before Tuesday @12noon	Q&A Session Tuesday @2pm	
		Video	Introduction	Pre-recorded		Tutorial Wednesday @1pm	
		Video	Background Models	Pre-recorded	Video before Tuesday @12noon		Assignment 1, Part 4
6	17/10/2022	Edges	Edge detection	Pre-recorded			
		Edges	Contour segmentation	Pre-recorded		Q&A Session Tuesday @2pm	
		Edges	Hough transform	Pre-recorded		Tutorial Wednesday @1pm	
7	24/10/2022	Edges	Least squared error & RANSAC	Pre-recorded	Edges before Tuesday @12noon		Assignment 1, Part 5
		Learning & Evaluation	Learning & Ground Truth	Pre-recorded			
		Learning & Evaluation	Performance	Pre-recorded	Learning & Evaluation before Tuesday @12noon	Q&A Session Tuesday @2pm	
8	31/10/2022	Recognition	SPR	Pre-recorded	Recognition before Tuesday @12noon	Tutorial Wednesday @1pm	
		Recognition	SPR Features	Pre-recorded			
		Regions II	Watershed segmentation	Pre-recorded			
9	07/11/2022	Regions II	Mean shift	Pre-recorded	Regions II before Tuesday @12noon	Q&A Session Tuesday @2pm	
		Recognition II	Template Matching	Pre-recorded		Tutorial Wednesday @1pm	
		Recognition II	Classifier Matching	Pre-recorded	Recognition II before Tuesday @12noon		Submit Assignment 1 - before 2-Nov @11:59pm
10	14/11/2022	Recognition III	SVM	Pre-recorded			Assignment 2 (10%)
		Recognition III	Haar	Pre-recorded		Q&A Session Tuesday @2pm	
		Recognition III	PCA	Pre-recorded	Recognition III before Tuesday @12noon	Tutorial Wednesday @1pm	- Sample Exam Question - released on 9-Nov @5:30pm
11	21/11/2022	Video II	Tracking - Mean Shift	Pre-recorded			
		Video II	Tracking - Optical Flow	Pre-recorded	Video II before Tuesday @12noon	Q&A Session Tuesday @2pm	Submit Assignment 3 - before 20-Nov @11:59pm
		Features	Corner detection - Moravec	Pre-recorded		Tutorial Wednesday @1pm	
12	28/11/2022	Features	Corner detection - Harris	Pre-recorded		Q&A Session Tuesday @2pm	
		Features	SIFT Keypoints	Pre-recorded		Tutorial Wednesday @1pm	
		Features	SIFT Matching & Recognition	Pre-recorded	Features before Tuesday @12noon		
		Conclusions	How to address vision problems	Live		Lecture Tuesday @2pm	
		Conclusions	Sample exam paper tutorial	Live		Tutorial Wednesday @1pm	

Introduction


Based on *A Practical Introduction to Computer Vision with
OpenCV* by Kenneth Dawson-Howe © Wiley & Sons Inc. 2014

Slide 2

2

Introduction – Human Vision

- ☞ When it was discovered that an image is projected inside the eye...
 - Scientists thought vision was solved.
 - But there is no homunculus in our brain.
- ☞ So how do we understand the images from the eyes?
 - 20% of the brain does “visual only” processing.
 - 40%+ of the brain does “vision+” processing.
 - Motor, attention, navigation, ...
- ☞ What does the brain do?
 - Visual cortex – basic image (pre)processing
 - Dedicated areas – face recognition, motion analysis, ...
 - We know very little really.



CC
SOME RIGHTS RESERVED
Courtesy of Pbroks13 and Jennifer Garcia, see Wikipedia article on Cartesian theatre for more information

Introduction

Based on A Practical Introduction to Computer Vision with
OpenCV by Kenneth Dawson-Howe © Wiley & Sons Inc. 2014

11:37 Slide 3

3

Introduction – Emulating human vision

- ☞ Machines can easily solve many tasks that we find difficult.
 - Particularly if we can describe them mathematically or in terms of an algorithm
- ☞ The solutions of many tasks that we find easy seem very difficult to describe formally.
 - Object recognition.
 - Rigid objects (variety)
 - Deformable objects
 - World modelling.
- ☞ Is machine vision just implementing human vision on a machine?
 - We don't understand human vision properly.
 - Machine vision is not restricted to two images from the visible spectrum.
 - UV, IR, Multiple cameras, Laser range finders, ...
 - Machine learning can “tackle tasks that are too difficult to solve with fixed programs written and designed by human beings.” From *Deep Learning* by Goodfellow



Introduction

Based on A Practical Introduction to Computer Vision with
OpenCV by Kenneth Dawson-Howe © Wiley & Sons Inc. 2014

11:37 Slide 4

4

Computer Vision is hard

"What we experience, apparently directly, is actually very different from what is recorded by our sense organs."
[Perception: From Sense to Object by J. Wilding]


```
9ca197af64736d708a82808c81948f8e9488717b6d9e98a2989bd1717b7a876d
9b99a4ad687a6e778290878a8380818885818b816895989f9a9bd5647d7d864e
9da4aea766756a6e79908e8c9681798784707b7e7587a3a19da38cda72845d35
a3b495a160656a74787a807d809ea3c580868d7c7785969a989697cf98502f35
ab969b9a5d626c7e807f8f74869ca9b4c9d37b7364819a9b958e978fd5393c38
9e619cb361616e727b7f6b86857aadadc7d2d860698e9a6f908d8ea33635365b
7556959a63686a77816a74781989faac6c6c3e5548c9c497992al2b4740479b
635f96ad5e5673729776828a898cbfc0d3d3c1c6e67b9c5527d3773c3f428fa6
626198a655686b77b075767f6eb1adbfbec7c2cccd7c9a3dacd1da3b2b35a3a1
697196b760615d6e9f6f79788aa3aeaca8bbb9c8c4c4d7d8c7c5823f36a3a2a0
6e5e9db25c606865ab7e5f84a58fabaaa29db7a2a7d4c2c3d282b398f9c9ba2
686balac5d636754b38d749e8f8f839a838a359ab8b9c7c154d732378fa5a09d
6a6996ae66606d6cbe8f827e8c823e763b2fcd0b8ca6592bc32c70a4a8a6a1
6e5f96ae645963705d72847a424331505abc2a7c2955fccc343c879ea5a293
6c6e96ae6862697d657a8143324434229fbfb1be88c83350323538a6979e999b
666b99b664666479867d60413f342ca0c794afbfc9d53b8b3d3788a6a498988c
626c98af686e6d809c60316f6d3283b9646d7ca3ca4848b3453e8ca3a59aa38d
6a6594b76d696fac37356639283fc47e69c7849a83ad37b13630a299999c8e8a
6b6088c1736f7832863b843635bb5f9aaeae898cb57f34a22f90al9a9982acc9
50548fb26e7232318e36328f68a7678ba6b2908e9f893e6c3f969a9a9584bbd2
5d498bb7756d709b9b382d2bb133628a9ca27876b18c3e2f24a29b9c284beddcb
56439db66c6b41406c68354f34a6d8a8a9997b3953141386c96989a83c5c0d0
3d3e9ab36d9a29447134d42f3e534d8a95707a85993962389c98928784d3cbce
383888b274642e6b7c4cb53363b456d89868595393f5957a09b96987ecbd8d0
31348dbd6f3d4c2d66852b4035313c3b3789b28e2e376a36998e9d8b7c1d1d5b8
813c7ab370262654366a93f303a366a7b9597c4cb606a4897816b6e93d49f4d
4a9b73b0a932243e6d647b693e3e4a5f8a869bb6c0d4834a978c9073cdd05d7c
2e9e72afa2363227307a71a73c3c6764898e8fa2abcad322988b77bade746a65
27b75dbcal303848265c652f5737806e91849a9fb9cbcf313999280cdd27c6c6c
28b759b8823f4c4b308c412c2b5c7e85928e9494abbbdcde56698b8fc154746f
```

Introduction

Computer Vision is about understanding images.

These images can be

- Greyscale or Colour or Multi-spectral
- Snapshots or video sequences
- Taken with a static or moving camera
- Taken of a stationary or dynamic scene
- Taken with a calibrated or un-calibrated camera




What Computer Vision aims to do is to extract some useful information from these images for...

- Inspection purposes
- Analysis purposes
- Control purposes



Applications – Industrial Inspection



© Reproduced by permission of James Mahon



© Reproduced by permission of Omron Electronics LLC

Introduction

Based on *A Practical Introduction to Computer Vision with OpenCV* by Kenneth Dawson-Howe © Wiley & Sons Inc. 2014

11:37 Slide 7

7

Applications – Surveillance / Forensics

- ❖ In London there are 500,000+ cameras
- ❖ However, only 1 crime solved per 1,000 cameras (BBC, 2009)
- ❖ In July 2005 there was a major terrorist attack in London...
- ❖ Afterwards
 - 17,000 hours of video was studied
 - Over 1 year



© Image courtesy of Feelart / FreeDigitalPhotos.net

Introduction

Based on *A Practical Introduction to Computer Vision with OpenCV* by Kenneth Dawson-Howe © Wiley & Sons Inc. 2014

11:37 Slide 8

8

Applications – Biometrics



No need for a key. Just your face. Won't be fooled by photos.

Chui Smart Doorbell from <http://iothought.com/chui-smart-doorbell/>


Introduction

Based on *A Practical Introduction to Computer Vision with OpenCV* by Kenneth Dawson-Howe © Wiley & Sons Inc. 2014

11:37 Slide 9

9

Applications – Landmine detection



© Reproduced by permission of Zouheir Fawaz




Introduction

Based on *A Practical Introduction to Computer Vision with OpenCV* by Kenneth Dawson-Howe © Wiley & Sons Inc. 2014


11:37 Slide 10

10

Applications – Medical Imaging



Fetal Head Measurements (BPD, HC)



© Family Practice Notebook (Reproduced with permission)



© Etan J. Tal, released under Creative Commons 3.0

Introduction

Based on *A Practical Introduction to Computer Vision with OpenCV* by Kenneth Dawson-Howe © Wiley & Sons Inc. 2014

Slide 11

11

Vision is confusing



"Every image is the image of a thing merely for him who knows how to read it, and who is enabled by the aid of the image to form an idea of the thing."

[Handbook of Physiological Optics
by H. Helmholtz]

Introduction

Based on *A Practical Introduction to Computer Vision with OpenCV* by Kenneth Dawson-Howe © Wiley & Sons Inc. 2014

11:37 Slide 12

12

Interfacing with the world



Ambient light sensor
Proximity sensor
Flood Illuminator
Infrared camera
Speaker
Microphone
7MP camera
Dot projector



© WAYRAY (Reproduced with permission)

Computer Vision

13

13

Applications – Augmenting reality




© Warner Bros. (Reproduced with permission)

Computer Vision

14

14

Applications – Intelligent Devices




Idea: Stream Light Between Snowflakes

Goal: High Light Throughput and Accuracy

Illustration adapted from de Charette (ICCP, 2012)

© Takeo Kanade, CMU



© Takeo Kanade, CMU

Computer Vision

15

15

Applications – Self drive car



© Takeo Kanade, Robotics Institute, Carnegie Mellon University (Reproduced with permission)

Introduction

Based on *A Practical Introduction to Computer Vision with OpenCV 2* by Kenneth Dawson-Howe © Wiley & Sons Inc. 2014

11:37 Slide 16

16

The Ultimate Goal

"...apprehension by the senses supplies after all, directly or indirectly, the material of all human knowledge, or at least the stimulus necessary to develop every inborn faculty of the mind. It supplies the basis for the whole action of man upon the outer world... For there is little hope that he who does not begin at the beginning of knowledge will ever arrive at it's end"

[The Recent Progress of the Theory of Vision by H. v. Helmholtz (1873)]

"If our long-sought quest to create autonomous anthropomorphic automata is to succeed, we must first impart perceptual abilities to machines".

[A Guided Tour of Computer Vision, by V. Nalwa]

Introduction

Based on *A Practical Introduction to Computer Vision with OpenCV* by Kenneth Dawson-Howe © Wiley & Sons Inc. 2014

11:37 Slide 17

17

The goal of Computer Vision?

Ultimately emulating this...



So, how are we doing?



© NASA / JPL, reproduced with permission



© Honda Motor Co. Ltd.
(Reproduced with permission)

Computer Vision

18

18

Vision will...

- ✓ Solve harder problems with less constraints.
- ✓ Be a big part of User I/O
- ✓ Enhance security.
- ✓ Diagnose medical conditions reliably.
- ✓ Control vehicles automatically.
- ✓ Identify suspects through forensic analysis
- ✓ Provide more understanding about the environment.



From http://www.telepresenceoptions.com/images/Minority_Report.gif

Goals

- ✓ It is intended that the student should, at the end of the course...
 - Understand the broad subject of computer vision.
 - Understand the main algorithmic processes used to manipulate images.
 - Understand how information may be extracted from images, and the associated problems.
 - Be able to describe the various operations both algorithmically and mathematically.
 - Be able to code and test basic vision algorithms.
 - Be able to develop potential solutions to complex vision problems.

Teaching Philosophy

- Flipped classroom
 - Prerecorded lectures to be studied in advance of class.
- Teaching by example.
- Hands-on experience of computer vision operations.
- Hands-on experience of coding computer vision operations
- Application of techniques worked on in group tutorials.
- Based on a single text +
 - RANSAC + PCA + SVM + K-means + Watershed +
 - Deep Learning + Stereo +...
- Lectures, slides & sample code provided online
 - On Blackboard (mymodule.tcd.ie)



Introduction

Based on *A Practical Introduction to Computer Vision with OpenCV* by Kenneth Dawson-Howe © Wiley & Sons Inc. 2014

Slide 21

21

Module delivery (most weeks)

- Around 4 pre-recorded lecture sessions (15-25 minutes each)
- 1 or 2 mini-tests to assess basic understanding of the pre-recorded material
 - Must be done before the Q&A session (or 0 marks).
- Q&A & 'advanced' topics session to discuss any issues/questions relating to the pre-recorded lecture material
 - Students MUST have reviewed the pre-recorded lecture sessions. If you have not reviewed them then please do not ask any questions during the Q&A session.
 - Time permitting we will look at some more advanced/applied topics in Vision such as Ethics, Text recognition, Lane Detection, Recognition using Deep Learning, Face Recognition, Head pose estimation, etc.
- Tutorial session to consider how to address real world problems using vision.
 - Students will be broken into small groups to come up with solutions (for 10-15 minutes), the course lecturer will visit the groups, and then when we come back together groups will present their solutions which will be discussed.
 - If the Q&A session does not have much to discuss, we may start the tutorial session during that hour.
- Discussion board

Introduction

Based on *A Practical Introduction to Computer Vision with OpenCV* by Kenneth Dawson-Howe © Wiley & Sons Inc. 2014

Slide 22

22

Module Assessment

- 10% Mini-tests which assess basic understanding of the pre-recorded material.
 - Must be done before the Q&A session (or 0 marks will be awarded).
 - Single attempt only.
- 30% Programming assignment.
 - Solving a real world vision problem and assessing how you have done.
- 10% Sample exam question(s) assignment.
 - Exam practice with feedback!
- 50% In person exam
 - 2 hours.
 - Some application of vision to real world problems type questions.
 - Some compare & contrast type questions.

Introduction

Based on *A Practical Introduction to Computer Vision with OpenCV* by Kenneth Dawson-Howe © Wiley & Sons Inc. 2014

Slide 23

23

Module contents

Week	Beginning	Folder	Topic	Format	Mini Tests (10%)	Live Classes	Practical Component
1	12/09/2022	Introduction	What is Computer Vision	Live		Lecture Tuesday @2pm	
		Introduction	Module Organisation	Live		Lecture Wednesday @1pm	Install OpenCV & TIPS.
2	19/09/2022	Images	Camera Model & Digitisation	Pre-recorded			Experiment with TIPS.
		Images	Colour Image models	Pre-recorded			Assignment 1 (30%), Part 1
		Images	Noise & Smoothing	Pre-recorded	Images before Tuesday @12noon	Q&A Session Tuesday @2pm	
		Binary	Thresholding	Pre-recorded		Tutorial Wednesday @1pm	
		Binary	Cleaning Binary Images	Pre-recorded	Binary before Tuesday @12noon		
3	26/09/2022	Regions	Connectivity & Connected Component Analysis	Pre-recorded			Assignment 1, Part 2
		Regions	kMeans clustering	Pre-recorded	Regions before Tuesday @12noon	Q&A Session Tuesday @2pm	
		Geometric	Transformations	Pre-recorded		Tutorial Wednesday @1pm	
		Geometric	Interpolation	Pre-recorded	Geometric before Tuesday @12noon		
		Histograms	1D & 3D histograms	Pre-recorded			Assignment 1, Part 3
4	03/10/2022	Histograms	Equalisation, Comparison & Back Projection	Pre-recorded	Histograms before Tuesday @12noon	Q&A Session Tuesday @2pm	
		Video	Introduction	Pre-recorded		Tutorial Wednesday @1pm	
5	10/10/2022	Video	Background Models	Pre-recorded	Video before Tuesday @12noon		Assignment 1, Part 4
		Edges	Edge detection	Pre-recorded			
		Edges	Contour segmentation	Pre-recorded		Q&A Session Tuesday @2pm	
		Edges	Hough transform	Pre-recorded		Tutorial Wednesday @1pm	
		Edges	Least squared error & RANSAC	Pre-recorded	Edges before Tuesday @12noon		
6	17/10/2022	Learning & Evaluation	Learning & Ground Truth	Pre-recorded			Assignment 1, Part 5
		Learning & Evaluation	Performance	Pre-recorded	Learning & Evaluation before Tuesday @12noon	Q&A Session Tuesday @2pm	
		Recognition	SPR	Pre-recorded		Tutorial Wednesday @1pm	
		Recognition	SPR Features	Pre-recorded	Recognition before Tuesday @12noon		
8	31/10/2022	Regions II	Watershed segmentation	Pre-recorded			
		Regions II	Mean shift	Pre-recorded	Regions II before Tuesday @12noon	Q&A Session Tuesday @2pm	
		Recognition II	Template Matching	Pre-recorded		Tutorial Wednesday @1pm	
		Recognition II	Classifier Matching	Pre-recorded			Submit Assignment 1 - before 2-Nov @11:59pm
		Recognition II	SVM	Pre-recorded	Recognition II before Tuesday @12noon		Assignment 2 (10%)
9	07/11/2022	Recognition III	Haar	Pre-recorded		Q&A Session Tuesday @2pm	
		Recognition III	PCA	Pre-recorded	Recognition III before Tuesday @12noon	Tutorial Wednesday @1pm	- Sample Exam Question - released on 9-Nov @5:30pm
10	14/11/2022	Video II	Tracking - Mean Shift	Pre-recorded			Submit Assignment 3 - before 20-Nov @11:59pm
		Video II	Tracking - Optical Flow	Pre-recorded	Video II before Tuesday @12noon	Q&A Session Tuesday @2pm	
11	21/11/2022	Features	Corner detection - Moravec	Pre-recorded			
		Features	Corner detection - Harris	Pre-recorded		Q&A Session Tuesday @2pm	
		Features	SIFT Keypoints	Pre-recorded		Tutorial Wednesday @1pm	
12	28/11/2022	Features	SIFT Matching & Recognition	Pre-recorded	Features before Tuesday @12noon		
		Conclusions	How to address vision problems	Live		Lecture Tuesday @2pm	
		Conclusions	Sample exam paper tutorial	Live		Tutorial Wednesday @1pm	

Introduction

Based on *A Practical Introduction to Computer Vision with OpenCV* by Kenneth Dawson-Howe © Wiley & Sons Inc. 2014

Slide 24

24